



Computer Communication Networks

Transport Layer


Dr. Raja Vara Prasad

Assistant Professor

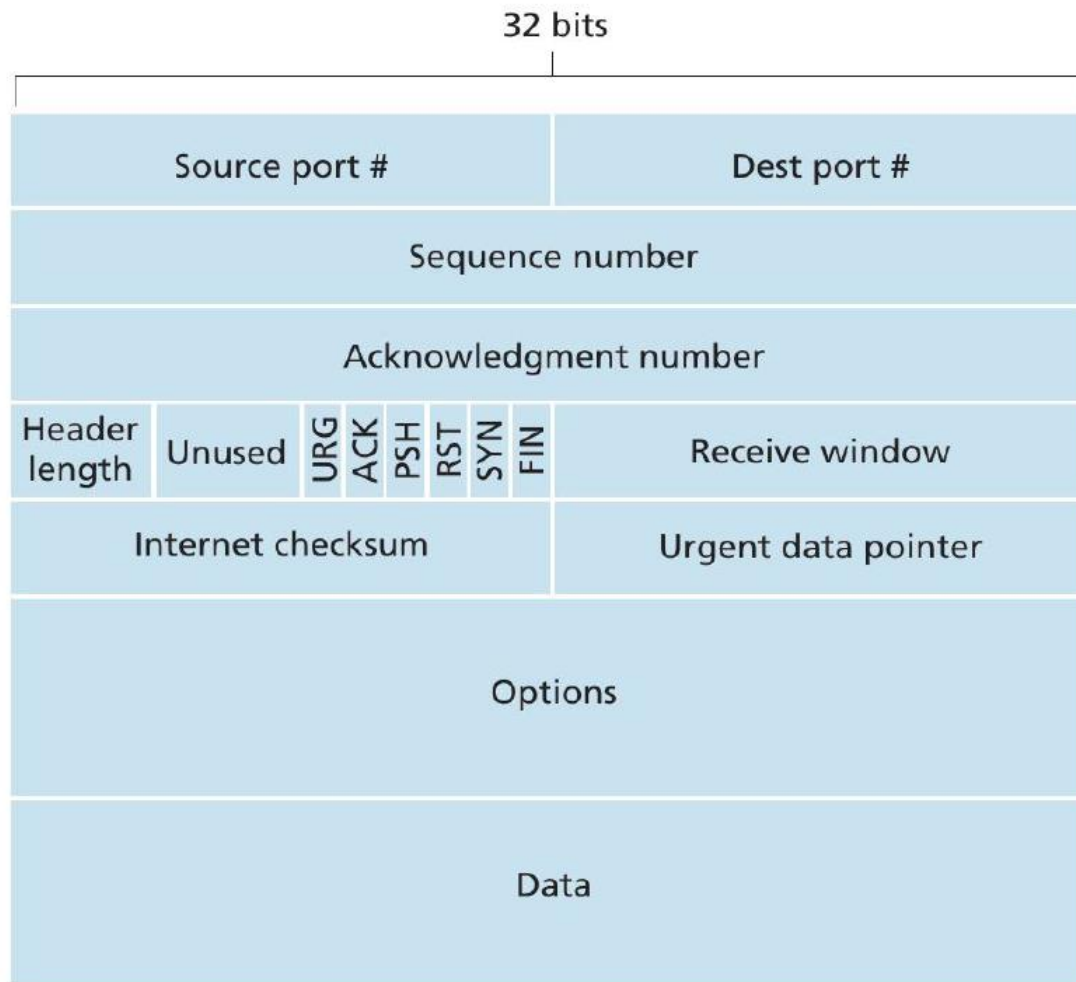
IIIT Sri City

TCP

- TCP is a **full duplex** service
- No **multicasting**
- **Maximum segment size (MSS)** is the maximum amount of **data** that a TCP segment can contain.



TCP Segment




TCP Segment

- The 16-bit **receive window** indicates the number of bytes that a receiver is willing to accept
- **Header length** field is 4-bytes, specifies the length of the TCP header in **32-bit words**.
- **Options** are used to negotiate MSS, include time-stamping, etc.
- The **flag field** contains 6 bits, **RST**, **SYN**, **FIN** are used for connection setup and teardown.
- **PSH** indicates that data has to be sent to upper layers immediately.
- **URG** is used to mark the segment as urgent, when it is on there will be a 16-bit **urgent data pointer field** at the end of urgent data.


TCP Sequence Numbers

- The sequence number of a segment is the **byte-stream number** of the first byte of data.
- The acknowledge number is the **sequence number of the next byte** that the receiver is expecting from source.
- TCP provides **cumulative acknowledgments**; **Out-of-order segments?**
- Sequence numbers may not always start from '0'.




TCP Timeout


- **SampleRTT**: RTT of a freshly transmitted packet. Computed for each RTT.
- Exponentially weighted moving average: **EstimatedRTT** = $(1 - \alpha)\text{EstimatedRTT} + \alpha \text{SampleRTT}$
- $\alpha = 0.125$
- **DevRTT** = $(1 - \beta) \text{DevRTT} + \beta | \text{SampleRTT} - \text{EstimatedRTT} |$
- $\beta = 0.25$
- Timeout = **EstimatedRTT** + 4. **DevRTT**




Example: TELNET




Example: Different scenarios of Reliability aspects



a



b



c

Doubling the Timeout interval:

intervals grow exponentially after each retransmission

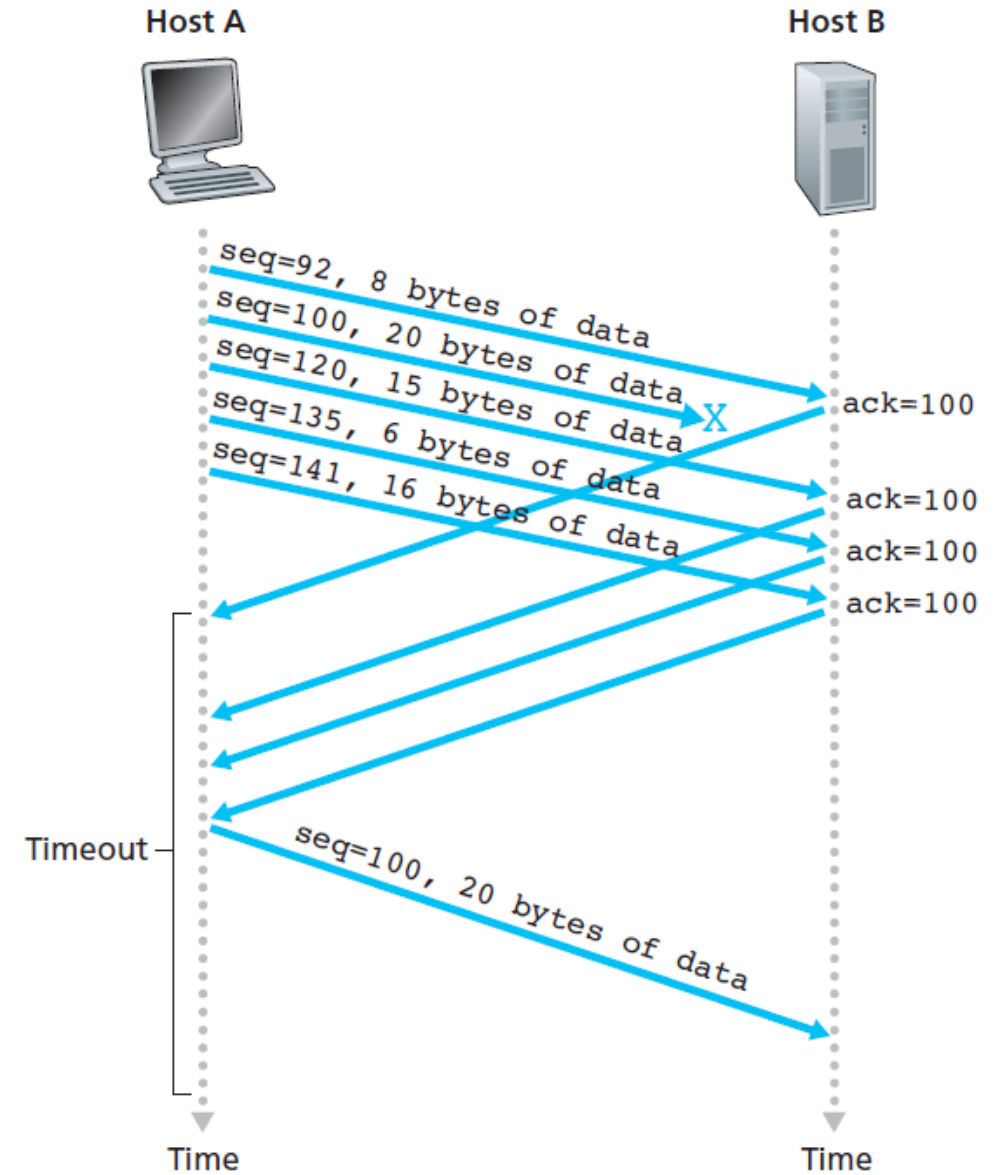
Fast Retransmit:

duplicate ACK multiple times

Example: three duplicate ACKs are received → the TCP sender performs a **fast retransmit**, retransmitting the missing segment *before* that segment's timer expires

Selection: Go-Back-N or Selective Repeat?

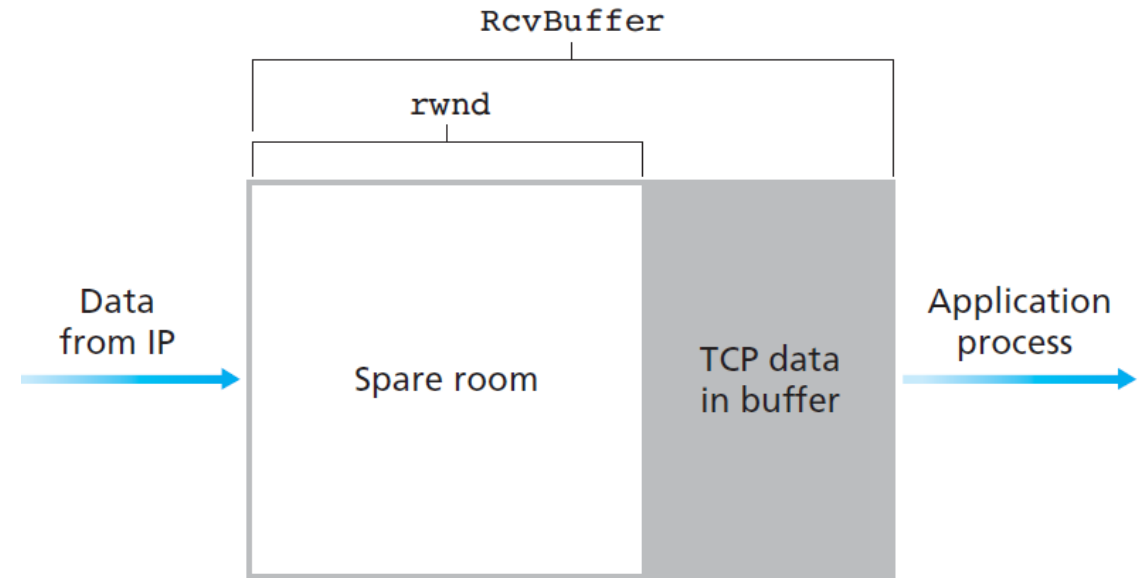
TCP's error-recovery mechanism is probably best categorized as a hybrid of GBN and SR protocols



Flow Control

Flow control is thus a speed-matching service → matching the rate at which the sender is sending against the rate at which the receiving application is reading

receive window: give the sender an idea of how much free buffer space is available at the receiver




$$\text{LastByteRcvd} - \text{LastByteRead} \leq \text{RcvBuffer}$$


The receive window, denoted **rwnd** is set to the amount of spare room in the buffer:

$$\text{rwnd} = \text{RcvBuffer} - [\text{LastByteRcvd} - \text{LastByteRead}]$$

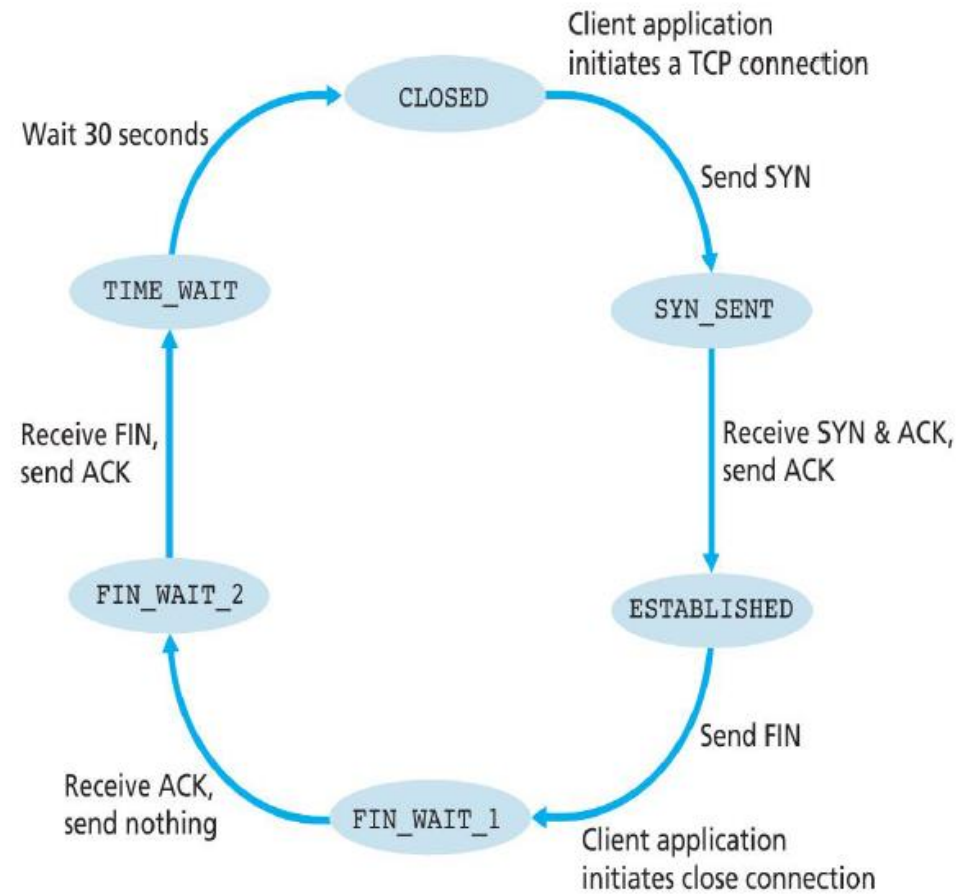
TCP Connection Establishment



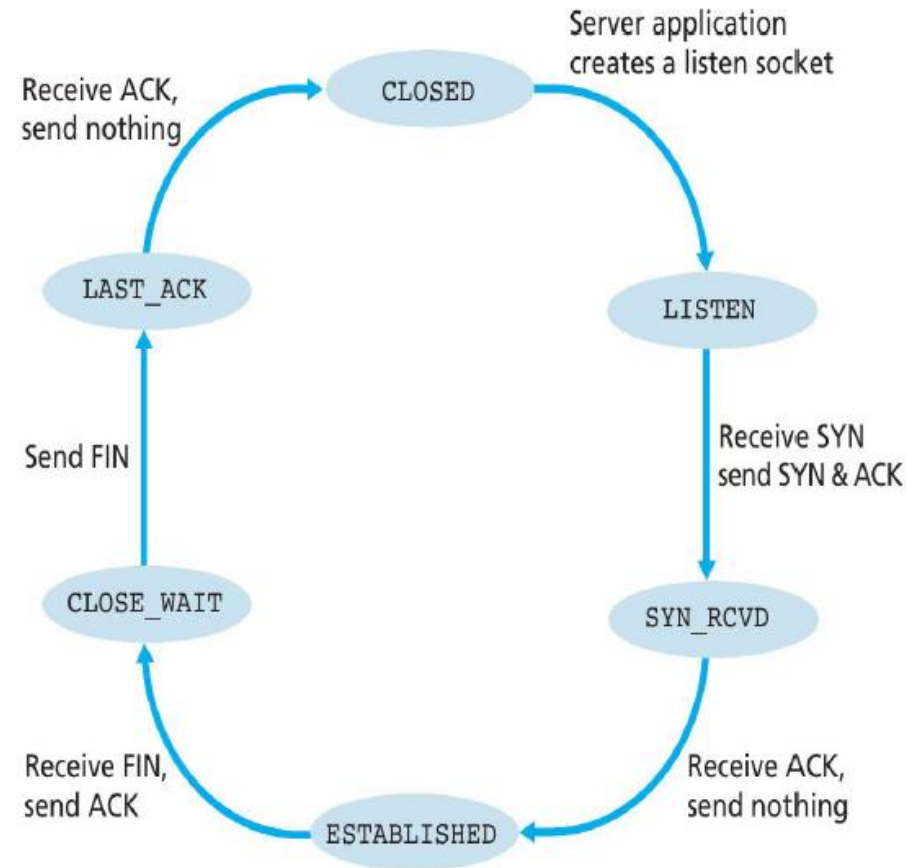
TCP Connection Termination



TCP States at Cleint




TCP States at Server



Congestion




- Why does Congestion occur?
- Packet arrival rate at a router is near or higher than the output link capacity.
- Consequences?
- Buffer overflows, retransmissions to compensate for lost packets
- Unneeded retransmissions

Congestion Scenario - 1





- (a) Host A knows whether buffer in the router has free space or not (Magic!)
- (b) Host A retransmits only if it is sure that packet is **lost** (Someone has to give this information)
- (c) Host A retransmits on timeouts!

Congestion Scenario - 1



Congestion Scenario - 2



Congestion Control

- End-to-end congestion control: no **explicit** information about congestion the network
- Network-assisted Congestion Control: **Choke packet**
- How does TCP identify congestion?
- No assistance from IP
- Identify congestion through **timeouts** and **duplicate ACKs**

Congestion Control





Figure 3.50 ♦ Congestion-control framework for ATM ABR service

EF CI bit. Each *data cell* contains an **explicit forward congestion indication (EF CI) bit**.


CI and NI bits: **congestion indication (CI) bit** and a **no increase (NI) bit**

ER setting. Each RM cell also contains a 2-byte **explicit rate (ER) field**.

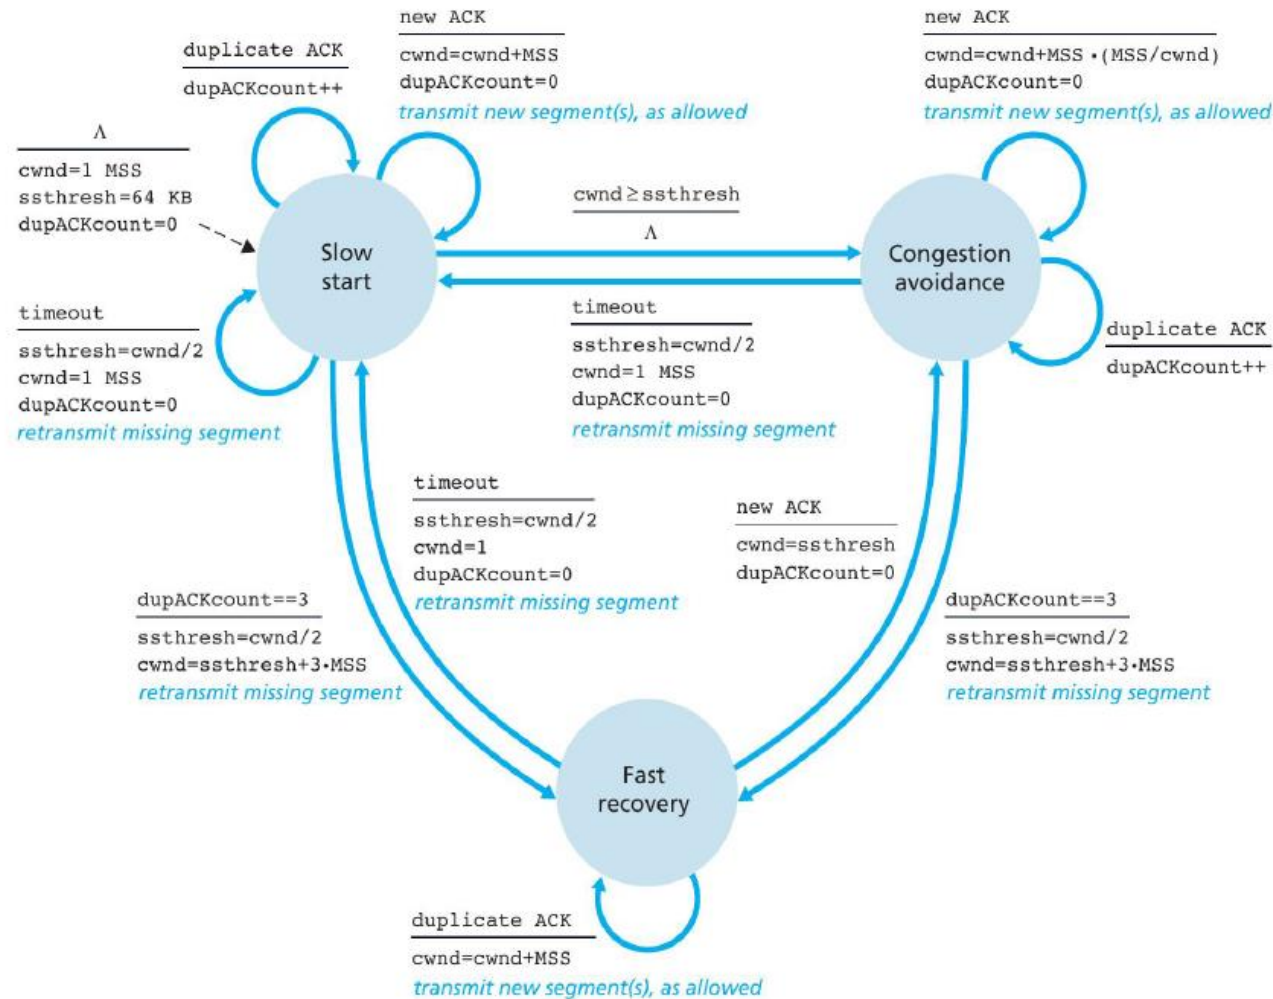
TCP's Congestion Control

- How does TCP control the sending rate?
- Defines a new variable called **cwnd**.
- $\text{LastByteSent} - \text{LastByteAcked} \leq \min\{\text{cwnd}, \text{rwnd}\}$
- Sends cwnd bytes of information per RTT (approximately)
- Can we adjust the speed? **Self-clocking**
 - A lost segment triggers the sender to reduce rate of transmission
 - An acknowledgment indicates **all is well!** Increase the rate
 - Bandwidth Probing


TCP Slow Start



TCP Congestion Control



TCP Congestion Window



AIMD




Figure 3.54 ♦ Additive-increase, multiplicative-decrease congestion control

additive-increase, multiplicative decrease (AIMD) form of congestion control:

TCP linearly increases its congestion window size (and hence its transmission rate) until a triple duplicate-ACK event occurs.

